



MediTrack

KUBERNETES CLUSTER

CONTENTS

Table of Figures.....	2
1 Solution.....	3
1.1 Introduction	3
1.2 Solution architecture - Minikube	3
1.3 Solution architecture – EKS	4
2 Microservices	5
2.1 Appointment Scheduling Service.....	5
2.2 Notification Service	6
2.3 Patient Records Service	6
2.4 Analytics and Aggregation Service	7
3 Deployment.....	9
3.1 Deployment method	9
4 Minikube cluster	10
5 MongoDB	13
6 Redshift.....	14
7 Quicksight Dashboard	15
8 CI/CD	16
9 Security and ethics	18
9.1.1 Security Challenges	18
9.1.2 Ethics Challenges	18
10 Runbook	20
10.1 Installing Minikube	20
10.2 Starting Minikube.....	20
10.3 Enabling Minikube Add-ons.....	21
10.4 Setting Up Microservices	21
10.5 Connecting MongoDB.....	22
10.6 Connecting to Amazon Redshift	22
10.7 Setting Up QuickSight	22
10.8 Configuring GitHub Actions.....	23
10.9 Testing the Workflow	23
Reference.....	24

TABLE OF FIGURES

Figure 1: Minikube Solution Architecture Diagram	3
Figure 2: AWS EKS Solution Architecture Diagram	4
Figure 3: Microservices Architecture Diagram	5
Figure 4: Appointments Scheduling Microservice	6
Figure 5: Notification Microservice.....	6
Figure 6: Patient Record Microservice	7
Figure 7: Analytics Microservice.....	7
Figure 8: Redshift Analytics Microservice	8
Figure 9: Deployment Diagram	9
Figure 10: Ingress Configuration and Folder Structure	10
Figure 11: Minikube Pods and Services	11
Figure 12: Minikube Pod Health Dashboard.....	11
Figure 13: Minikube Dashboard - Cron Jobs and Deployments	11
Figure 14: Minikube Dashboard Ingress.....	12
Figure 15: MongoDB	13
Figure 16: AWS Redshift Queries.....	14
Figure 17: AWS Redshift Query Results	14
Figure 18: AWS QuickSight Dashboard.....	15
Figure 19: Data Sources for Quicksight	15
Figure 20: Github Actions Red-Green Workflow	16
Figure 21: Github Action CI/CD	16
Figure 22: Github Folder Structure	17
Figure 23: Minikube Version Check	20
Figure 24: Minikube Status Check.....	20
Figure 25: Minikube Addon list	21
Figure 26: MongoDB URL	22
Figure 27: Redshift Python Driver in Requirements.txt.....	22

1 SOLUTION

1.1 INTRODUCTION

HealthSync, developed for MediTrack, streamlines patient data management and appointment tracking, enabling healthcare providers to efficiently access and update health records.

1.2 SOLUTION ARCHITECTURE - MINIKUBE

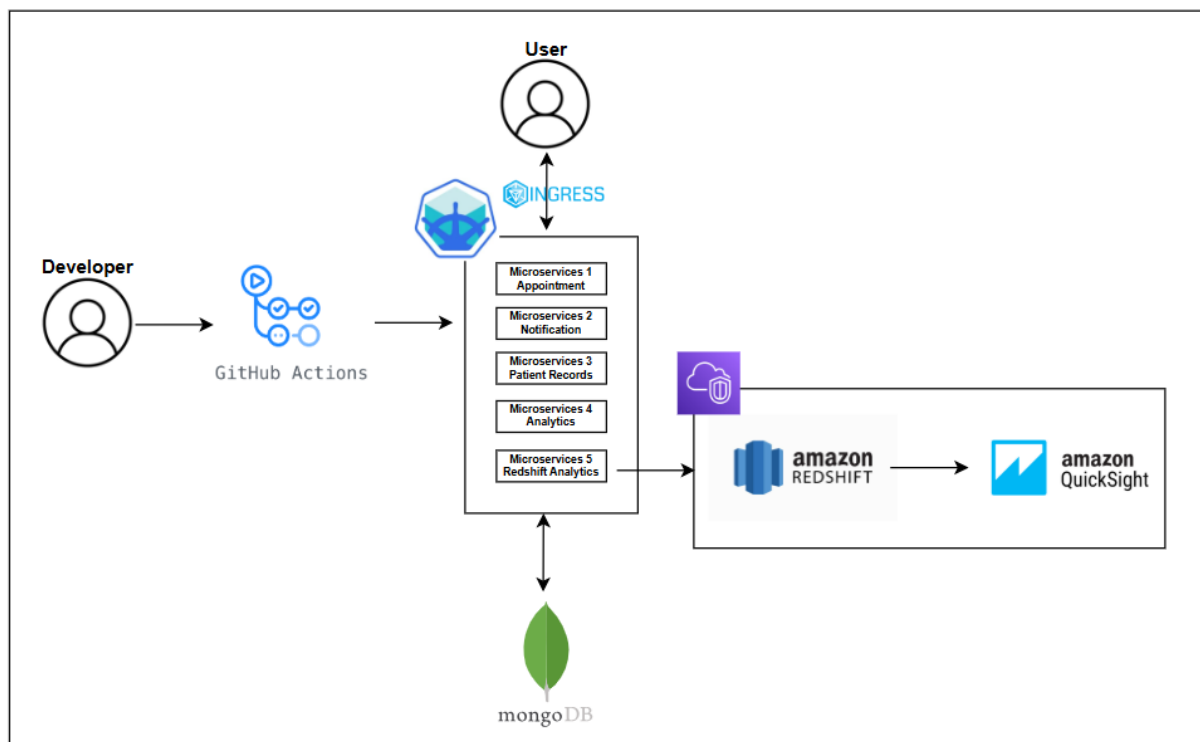


Figure 1: Minikube Solution Architecture Diagram

The solution is deployed on a local Kubernetes environment using Minikube. Minikube is configured to run on the local machine with Docker as the container runtime. Docker images are built directly on the Minikube virtual machine and deployed as Kubernetes pods. A Python Flask microservice handles incoming user requests and stores data in a locally hosted MongoDB instance. Another microservice consumes data from MongoDB, performs analytical processing, and sends the results to Amazon Redshift, which is hosted in an AWS VPC with appropriate security configurations. Visualizations of the analytical data are generated using Amazon QuickSight, which directly connects to the Redshift cluster. Continuous integration and deployment (CI/CD) are automated using GitHub Actions, which builds the Docker images, runs tests, and deploys updated services to the Minikube cluster.

1.3 SOLUTION ARCHITECTURE – EKS

The proposed solution migrates the entire infrastructure to AWS for improved scalability, security, and performance. The application is deployed in an Amazon VPC with both public and private subnets. The Kubernetes cluster is hosted on Amazon Elastic Kubernetes Service (EKS), which manages the orchestration of containerized microservices. Docker images are stored in Amazon Elastic Container Registry (ECR) and pulled by EKS during deployments. User requests are routed through an Internet Gateway to an Application Load Balancer (ALB), which directs traffic to the Flask microservice running within the EKS cluster. The Flask microservice interacts with MongoDB, which is either hosted on MongoDB Atlas or self-managed in a private subnet within the VPC. The analytics microservice processes data from MongoDB and pushes the results to Amazon Redshift, also hosted within the VPC. Amazon QuickSight connects to Redshift to generate visual dashboards and reports. CI/CD pipelines are implemented using GitHub Actions, automating the process of building Docker images, pushing them to ECR, and updating deployments in the EKS cluster.

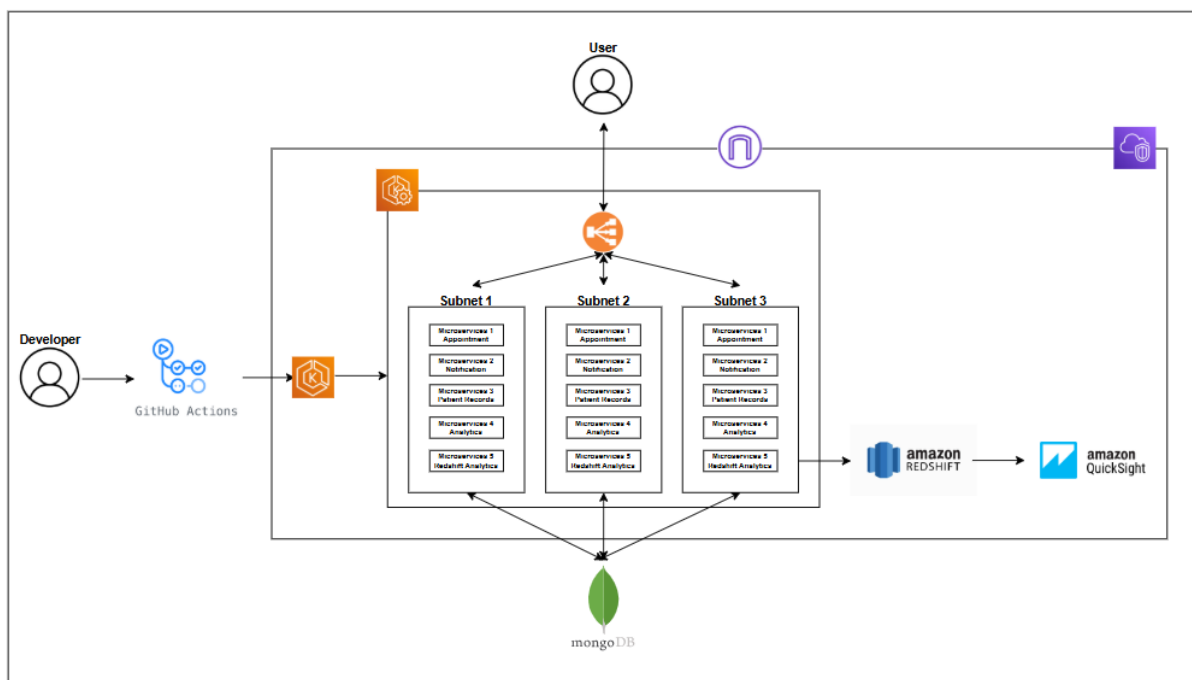


Figure 2: AWS EKS Solution Architecture Diagram

2 MICROSERVICES

I have implemented 4 microservices to efficiently handle data in this solution as,

- Appointment Scheduling Service
- Notification Service
- Patient Records Service
- Analytics and Aggregation Service and Redshift Analytics Service

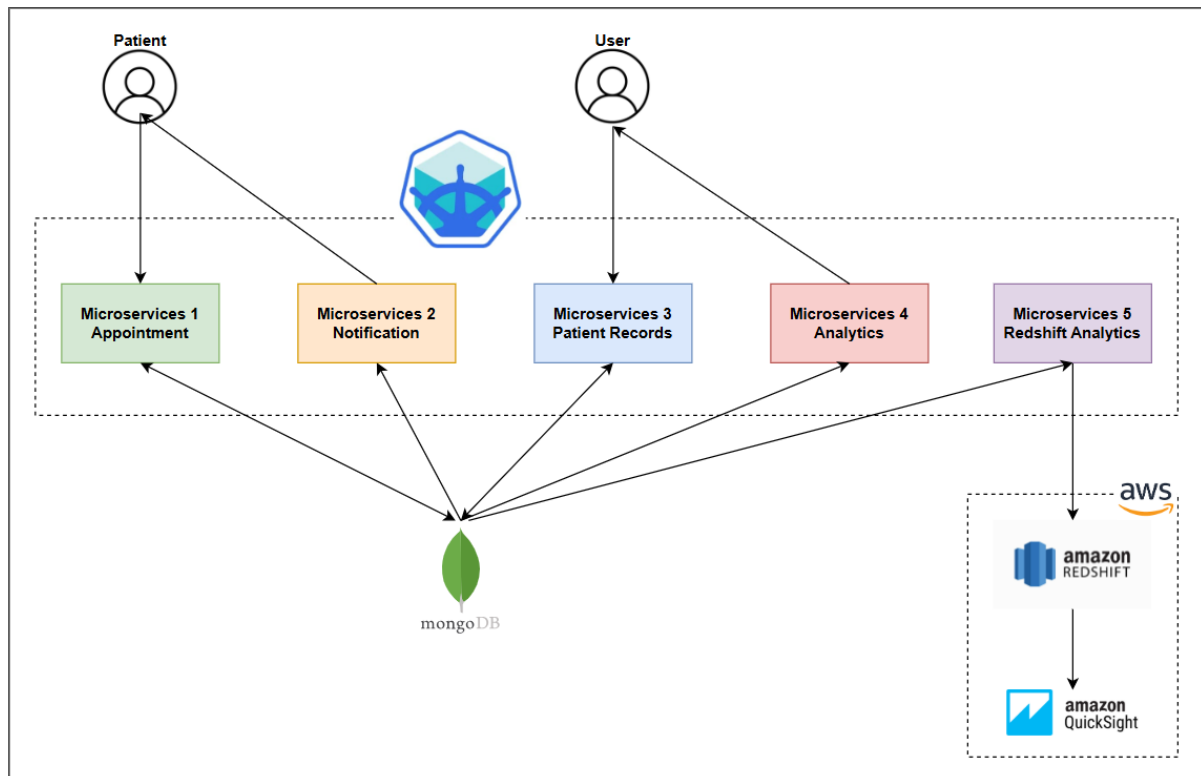


Figure 3: Microservices Architecture Diagram

2.1 APPOINTMENT SCHEDULING SERVICE

The Appointment Scheduling Service manages the core functionality of scheduling and maintaining appointments between patients and doctors. When a new appointment is scheduled, the service validates the incoming data for completeness and then saves it into the database. Upon successful insertion, it communicates with the Notification Service to send confirmation messages to patients, ensuring a seamless and connected experience. Key functionalities include,

- Adding Appointments
- Updating and Deleting Appointments
- Fetching Appointments

The service is designed to be robust and scalable, ensuring it meets the high demand of managing large volumes of appointment data efficiently.

```
}
tharindi@Tharindi:~/BDA/MediTrack/BDA-MediTrack$ curl http://meditrack.local/appointments
[
  {
    "_id": "6764686e25a3bc75b2ff56d3",
    "patient_name": "John Doe"
  },
  {
    "_id": "676468d625a3bc75b2ff56d4",
    "patient_name": "John Doe"
  },
  {
    "_id": "6764690625a3bc75b2ff56d5",
    "patient_name": "John Doe"
  },
  {
    "_id": "67646b1a25a3bc75b2ff56d6",
    "patient_name": "Jena Lary"
  },
  {
    "_id": "67646bd225a3bc75b2ff56d7",
    "patient_name": "Check new"
  },
  {
    "_id": "67646ddd25a3bc75b2ff56d8",
    "patient_name": "Tharindi Wi"
  }
]
```

Figure 4: Appointments Scheduling Microservice

2.2 NOTIFICATION SERVICE

The Notification Service complements the Appointment Scheduling Service by handling all patient communication related to appointments. This includes sending notifications about newly scheduled appointments and reminders for upcoming ones. It also queries the Appointments collection directly to identify upcoming appointments and send timely reminders. Key functionalities include,

- Scheduled Notifications
- Reminders sending
- Error Handling

The Notification Service adds value by ensuring patients are always informed about their healthcare schedules, improving reliability and trust in the system.

```
{ "error": "Patient not found" }
tharindi@Tharindi:~/BDA/MediTrack/BDA-MediTrack$ curl http://meditrack.local/notifications/reminders
{
  "details": [
    "Reminder: Hello testw, you have an appointment with Dr. Dr. Sam tomorrow at 10:00 AM."
  ],
  "message": "Reminders sent!"
}
tharindi@Tharindi:~/BDA/MediTrack/BDA-MediTrack$ curl -X POST -H "Content-Type: application/json" -d '{"patient name": "Ingress testnew", "doctor name": "Dr
```

Figure 5: Notification Microservice

2.3 PATIENT RECORDS SERVICE

The Patient Records Service is responsible for managing patient-related data such as demographics, medical history, prescriptions, and lab results. It stores data in the Patients collection and provides an interface to add, update, or delete records.

This service is critical for maintaining up-to-date patient information, which can be referenced by other microservices to provide personalized care and insights. Key functionalities include,

- Adding New Patients

- Updating Patient Information
- Deleting Patient Records

This service ensures data integrity and acts as a single source of truth for patient information across the system.

```
tharindi@Tharindi:~/BDA/MediTrack/BDA-MediTrack$ curl http://meditrack.local/patients
[[{"age":28,"condition":"Fever","name":"Jane Doe"}, {"age":30,"gender":"Female","medical history":["Asthma","High Blood Pressure"],"patient_name":"Jane Doe"}, {"condition":"Fever","date":"2024-12-20","doctor_name":"Dr. Smith","patient_name":"John Doe","time":"10:00 AM"}, {"condition":"Fever","date":"2024-12-21","doctor_name":"Dr. Smith","patient_name":"Tny Bell","time":"10:30 AM"}, {"condition":"Asthma","date":"2024-12-21","doctor_name":"Dr. Chi","patient_name":"Tharindi W","time":"11:50 AM"}, {"name":"Ingress test"}, {"age":25,"condition":"Cough","name":"Test Patient"}, {"name":"Ingress test"}]]
tharindi@Tharindi:~/BDA/MediTrack/BDA-MediTrack$ curl -X POST -H "Content-Type: application/json" -d '{"name": "Ingress testnew"}' http://meditrack.local/p
```

Figure 6: Patient Record Microservice

2.4 ANALYTICS AND AGGREGATION SERVICE

The Analytics and Aggregation Service provides insightful reporting and data analysis based on the information stored by other microservices. It performs scheduled jobs and uses MongoDB aggregation pipelines to derive actionable insights, which help improve healthcare operations and identify trends. Key functionalities include,

1. Number of Appointments per Doctor:
 $A_d = \text{Count (appointments per doctor)}$
2. Frequency of Appointments Over Time:
 $A_t = \text{Count (appointments per date)}$
3. Common Symptoms and Conditions by Specialty:

Groups medical conditions recorded during appointments by the specialty of the attending doctor.

4. Average Appointments per Patient:
$$Avg_p = \frac{\sum \text{appointments per patient}}{\text{total number of patients}}$$

These analytical reports enable hospitals and clinics to make data-driven decisions, optimize their services, and enhance patient outcomes. Initially this was tested with analytics microservice then using a new microservice as redshift analytics same results were pushed to AWS Redshift.

```
}
tharindi@Tharindi:~/BDA/MediTrack/BDA-MediTrack$ curl http://meditrack.local/aggregation/average_appointments_per_patient
{
  "average_appointments_per_patient": 1.875
}
tharindi@Tharindi:~/BDA/MediTrack/BDA-MediTrack$
```

Figure 7: Analytics Microservice


```

tharindi@Tharindi:~/BDA/MediTrack/BDA-MediTrack$ nc -zv myredshiftcluster.csqajvggkn5v.us-east-1.redshift.amazonaws.com 5439
Connection to myredshiftcluster.csqajvggkn5v.us-east-1.redshift.amazonaws.com (3.210.42.120) 5439 port [tcp/*] succeeded!
tharindi@Tharindi:~/BDA/MediTrack/BDA-MediTrack$ curl -X POST http://meditrack.local/sync/appointments_per_doctor
{
  "message": "Appointments per doctor synced to Redshift"
}
tharindi@Tharindi:~/BDA/MediTrack/BDA-MediTrack$ curl -X POST http://meditrack.local/sync/appointments_over_time
curl -X POST http://meditrack.local/sync/common_conditions_by_specialty
curl -X POST http://meditrack.local/sync/average_appointments_per_patient
{
  "message": "Appointments over time synced to Redshift"
}
{
  "message": "Common conditions by specialty synced to Redshift"
}
{
  "message": "Average appointments per patient synced to Redshift"
}
tharindi@Tharindi:~/BDA/MediTrack/BDA-MediTrack$

```

Ln 61, Col 1 (211 selected) Spaces: 2

Figure 8: Redshift Analytics Microservice

3 DEPLOYMENT

3.1 DEPLOYMENT METHOD

The deployment process begins with a CI/CD pipeline using GitHub Actions, which automates the build and deployment of Dockerized microservices on a Minikube cluster running locally. The pipeline initializes the Minikube environment, builds Docker images within the Minikube VM, applies Kubernetes manifests to deploy the microservices, and runs integration tests to verify functionality. The microservices include a Flask-based API for data storage in MongoDB Atlas and an analytics service scheduled via a Cron job to process data and push results to AWS Redshift daily. This workflow ensures consistent deployments and streamlined testing, with all updates managed through version control in GitHub.

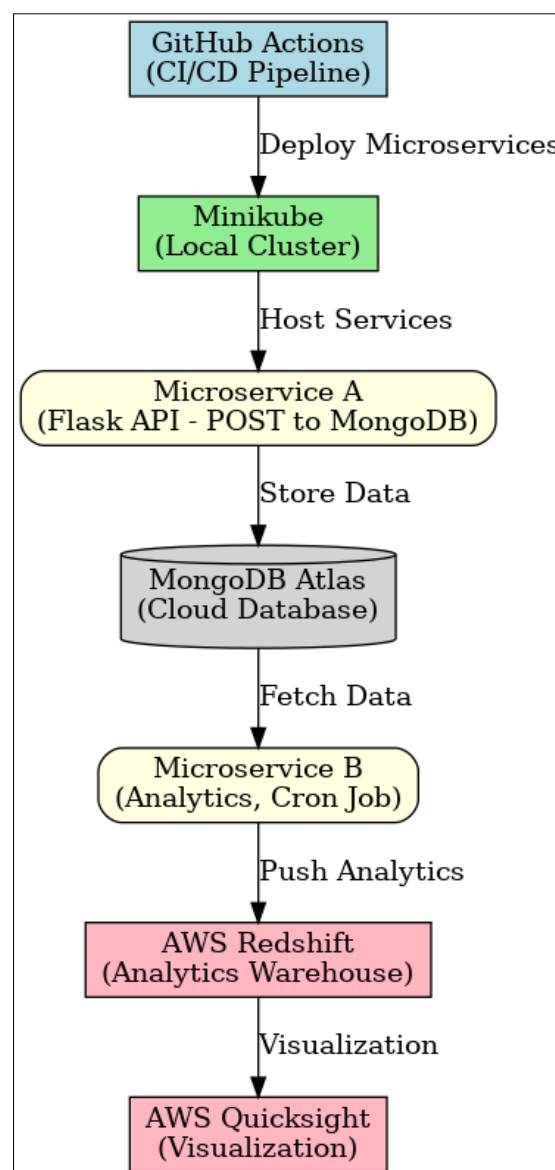


Figure 9: Deployment Diagram

4 MINIKUBE CLUSTER

The Minikube cluster hosts all microservices in a controlled local Kubernetes environment, with each microservice running in its own pod. An NGINX ingress controller, configured through the provided ingress manifest, manages load balancing and routing based on predefined rules, ensuring seamless API access via the meditrack.local domain. This setup directs requests to appropriate backend services like patient records, notifications, appointment scheduling, analytics, and Redshift data synchronization, using their respective paths and ports. Observability is enhanced using the Minikube dashboard.

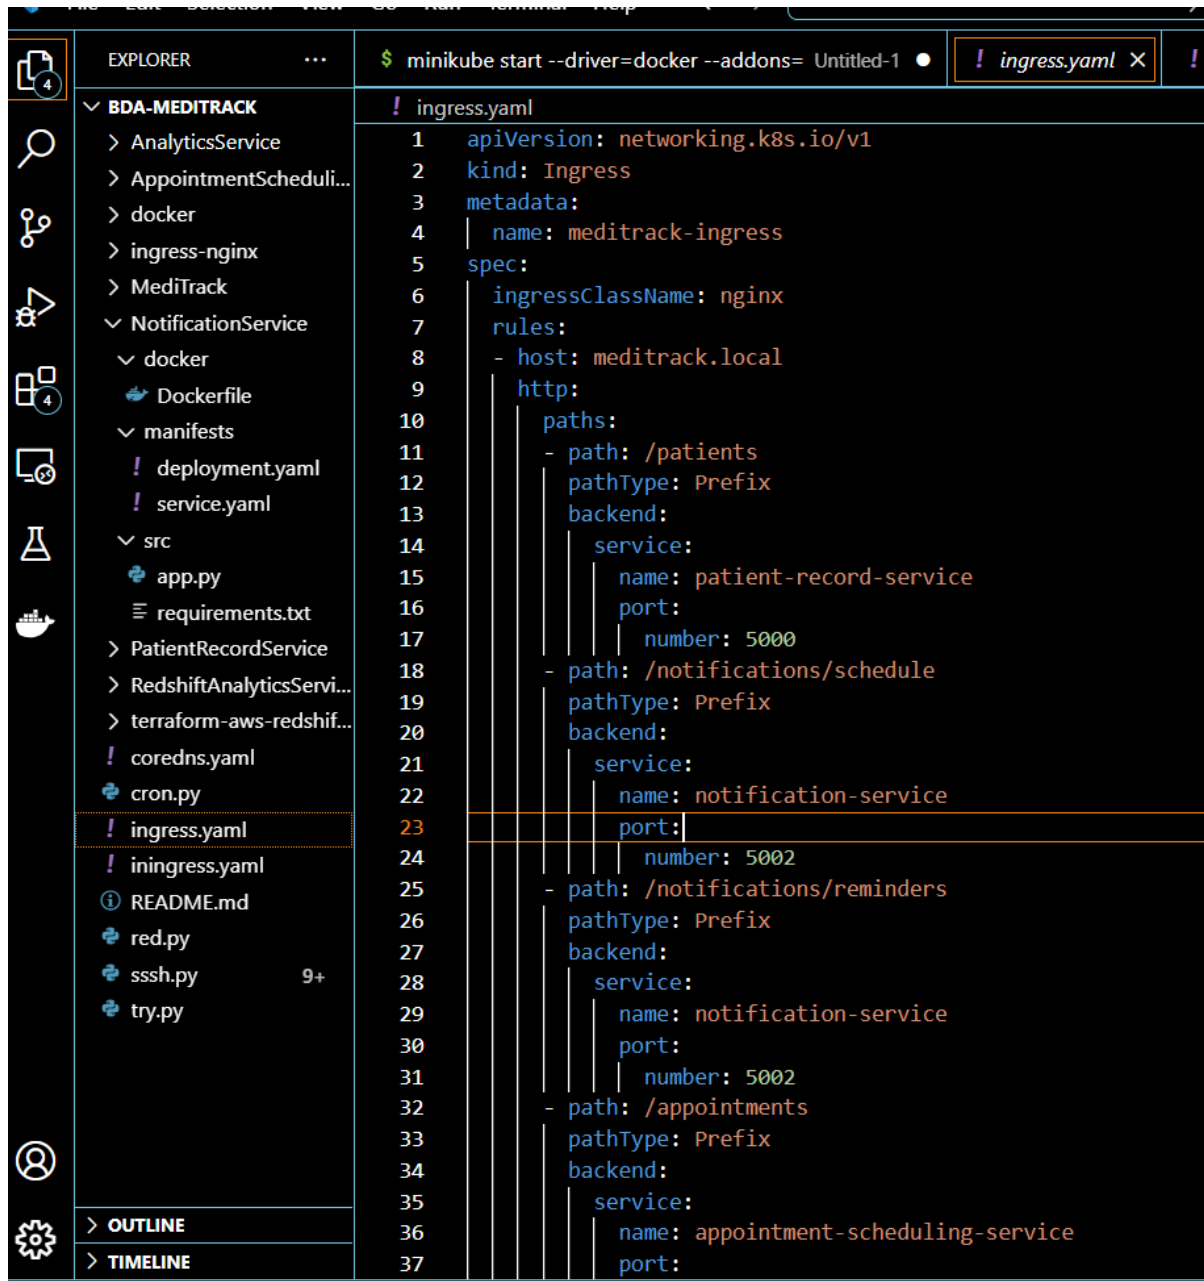


Figure 10: Ingress Configuration and Folder Structure

```

PROBLEMS 1K+ OUTPUT DEBUG CONSOLE TERMINAL PORTS
Password for 'https://tharindi-w@github.com':
Everything up-to-date
tharindi@tharindi:~/BDA/MediTrack/BDA-MediTrack/MediTrack$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
analytics-service-68844485b-ltqg6   1/1     Running   0           20h
appointment-scheduling-service-fd6cf786b-kzkhp 1/1     Running   0           20h
bar-app                              1/1     Running   0           20h
foo-app                              1/1     Running   0           20h
hello-minikube-7d48979fd6-w8swg     1/1     Running   0           20h
notification-service-5fcd84fddd-glpgw 1/1     Running   0           20h
patient-record-service-7b6cdb5bd4-22b7x 1/1     Running   0           20h
redshift-analytics-manual-job-w48pq   1/1     Running   0           19h
redshift-analytics-service-8hgdp     1/1     Running   0           19h
tharindi@tharindi:~/BDA/MediTrack/BDA-MediTrack/MediTrack$ kubectl get services
NAME                                TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)          AGE
analytics-service                   ClusterIP    10.96.250.204 <none>        5003/TCP         20h
appointment-scheduling-service       ClusterIP    10.99.175.187 <none>        5001/TCP         20h
bar-service                         ClusterIP    10.97.36.73   <none>        8080/TCP         20h
foo-service                         ClusterIP    10.107.21.237 <none>        8080/TCP         20h
hello-minikube                     NodePort     10.109.253.229 <none>        8080:30381/TCP   20h
kubernetes                         ClusterIP    10.96.0.1     <none>        443/TCP          20h
notification-service                ClusterIP    10.100.107.57 <none>        5002/TCP         20h
patient-record-service              ClusterIP    10.104.78.117 <none>        5000/TCP         20h
redshift-analytics-service           ClusterIP    10.103.170.87 <none>        5005/TCP         20h
tharindi@tharindi:~/BDA/MediTrack/BDA-MediTrack/MediTrack$ curl http://meditrack.local/aggregation/appointments.per.doctor

```

Figure 11: Minikube Pods and Services

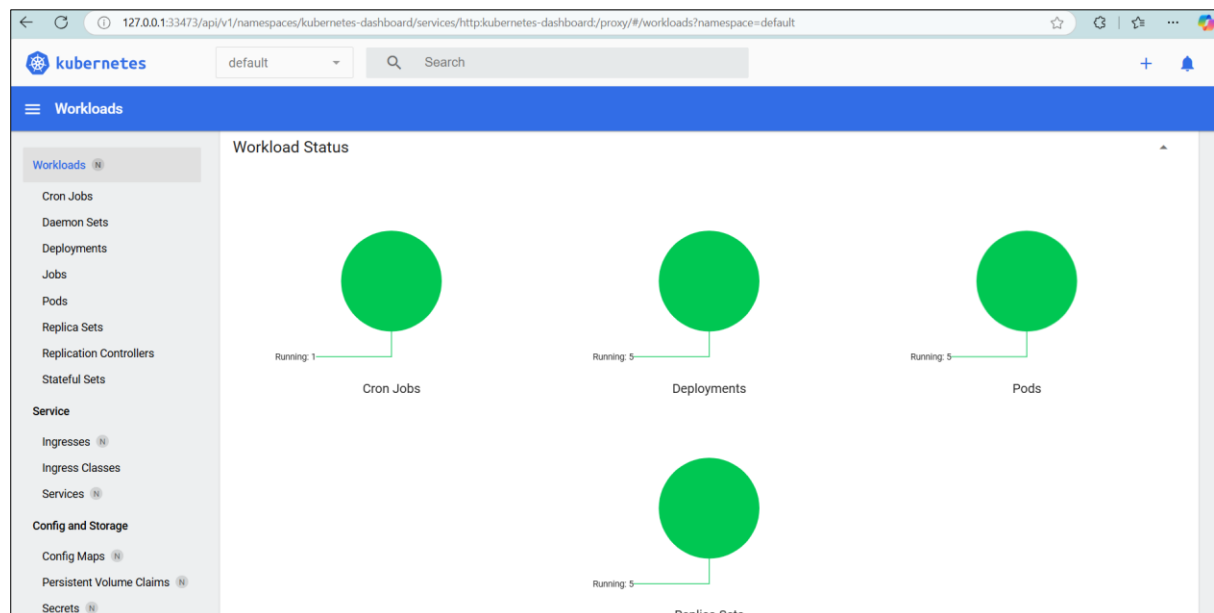
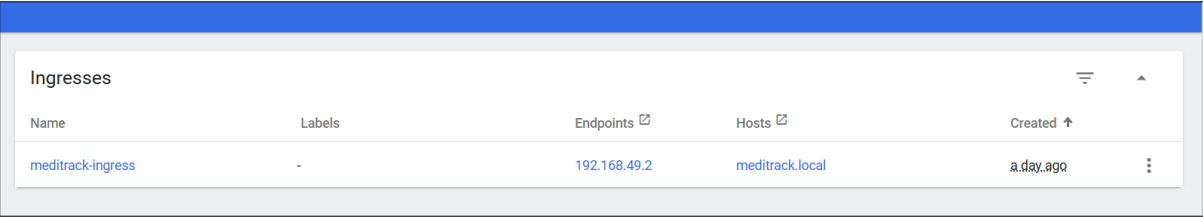


Figure 12: Minikube Pod Health Dashboard

Cron Jobs								
Name	Images	Labels	Schedule	Suspend	Active	Last Schedule	Created ↑	
redshift-analytics-cronjob	redshift-analytics:latest	-	0 4 ***	false	0	-	26.minutes.ago	

Deployments				
Name	Images	Labels	Pods	Created ↑
redshift-analytics	redshift-analytics:latest	-	1 / 1	a.day.ago
appointment-scheduling-service	appointment-scheduling-service:latest	-	1 / 1	2.days.ago
notification-service	notification-service:latest	-	1 / 1	2.days.ago
analytics-service	analytics-service:latest	-	1 / 1	2.days.ago
patient-record-service	patient-record-service:latest	-	1 / 1	2.days.ago

Figure 13: Minikube Dashboard - Cron Jobs and Deployments



The screenshot shows the 'Ingresses' section of the Minikube Dashboard. It features a table with the following columns: Name, Labels, Endpoints (with an external link icon), Hosts (with an external link icon), and Created (with an upward arrow icon). A single row is displayed with the name 'meditrack-ingress', no labels, the endpoint '192.168.49.2', the host 'meditrack.local', and a creation time of 'a.day.ago'. A vertical ellipsis menu icon is visible at the end of the row.





Name	Labels	Endpoints 	Hosts 	Created 	
meditrack-ingress	-	192.168.49.2	meditrack.local	a.day.ago	

Figure 14: Minikube Dashboard Ingress

5 MONGODB

MongoDB Atlas, a managed NoSQL database, was chosen for its scalability, flexibility, and ability to handle semi-structured medical data efficiently. Its distributed architecture ensures high availability and performance. Flask-based microservices connect securely to the cluster using a connection string and the pymongo driver. This setup ensures seamless integration, secure communication, and efficient data management.

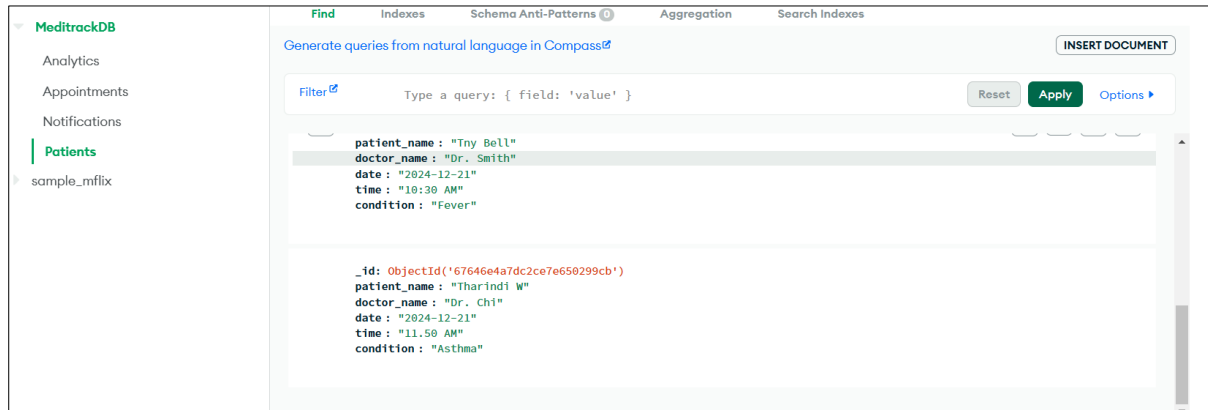


Figure 15: MongoDB

6 REDSHIFT

AWS Redshift is deployed within a VPC, leveraging three subnets across different availability zones for high availability and fault tolerance. Custom route tables are configured for secure data flow within the VPC, while IAM roles enable granular access control for data ingestion and analytics tasks. Security groups restrict inbound and outbound traffic, ensuring only authorized services and Quicksight have access to the database. This setup provides a robust and secure data warehouse for processing analytics.

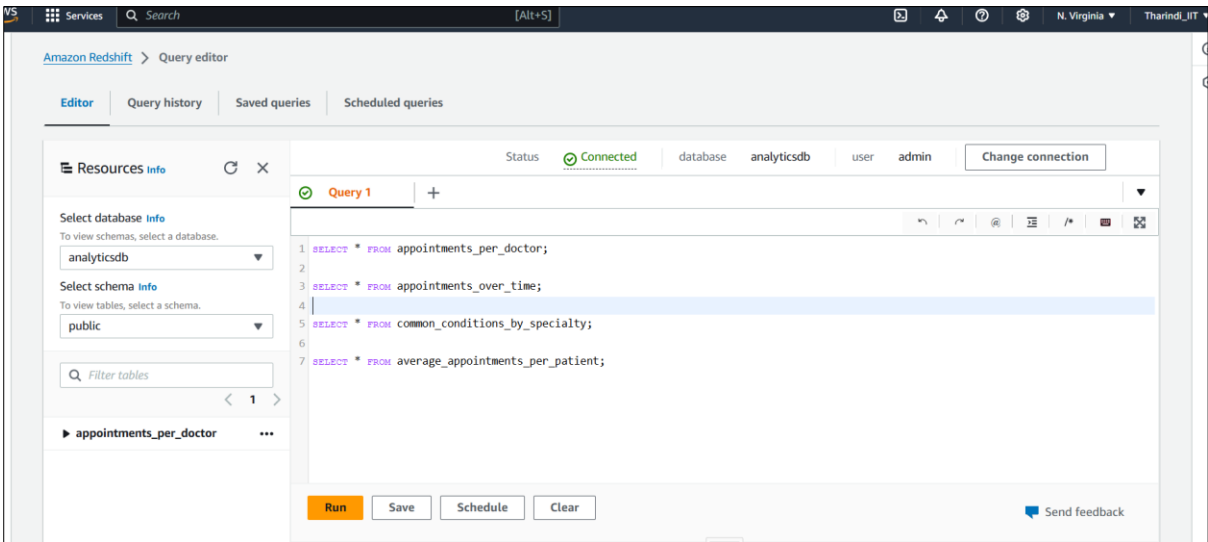


Figure 16: AWS Redshift Queries

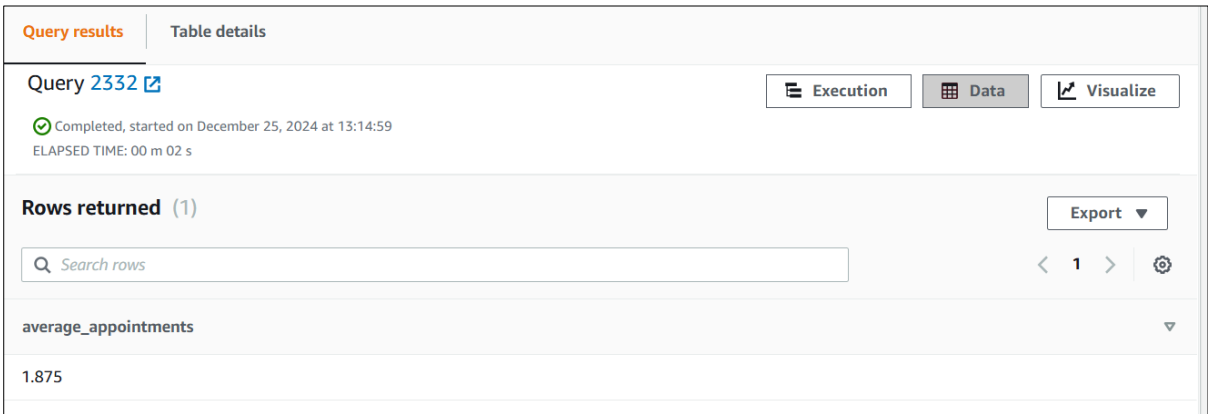


Figure 17: AWS Redshift Query Results

7 QUICKSIGHT DASHBOARD

AWS QuickSight integrates with Redshift to deliver interactive visualizations and dashboards. QuickSight's user-friendly interface and rich visualization capabilities

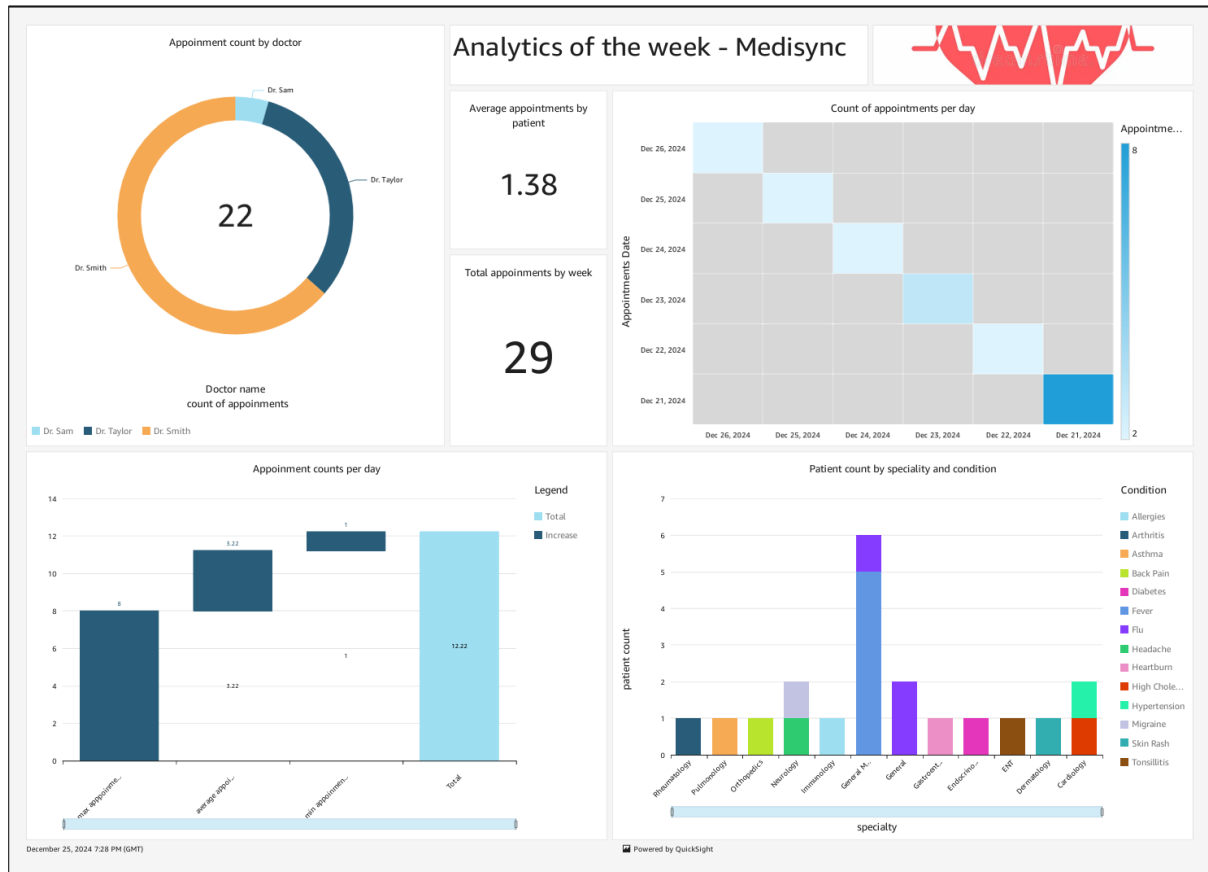


Figure 18: AWS QuickSight Dashboard

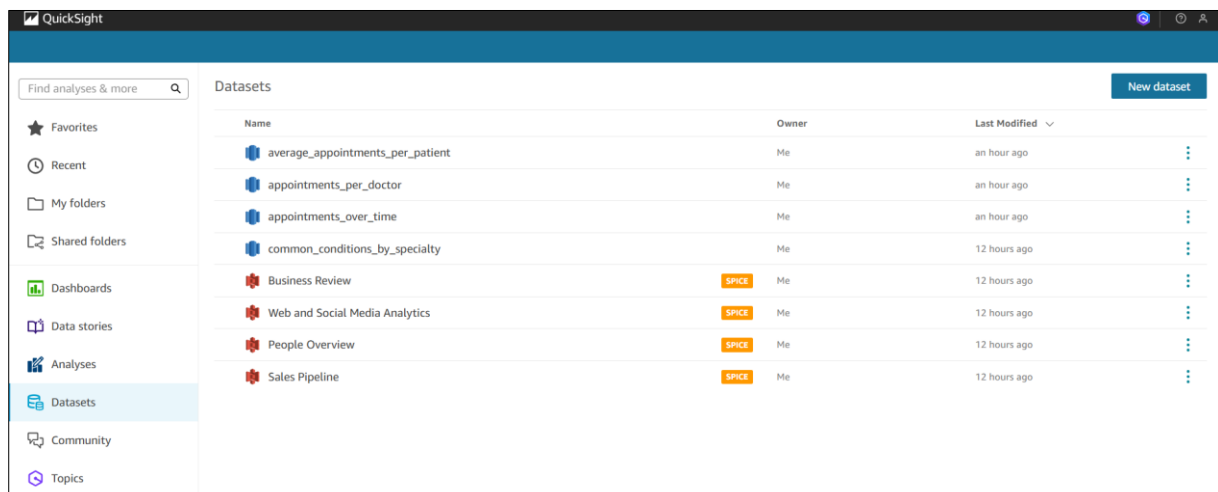


Figure 19: Data Sources for QuickSight

8 CI/CD

The CI/CD pipeline, implemented using GitHub Actions, automates the deployment process for microservices on a Minikube cluster. It initializes Minikube, builds Docker images in the Minikube VM, applies Kubernetes manifests, and performs integration tests to ensure service functionality. This setup ensures seamless version control, automated updates, and reliable testing before production deployment.

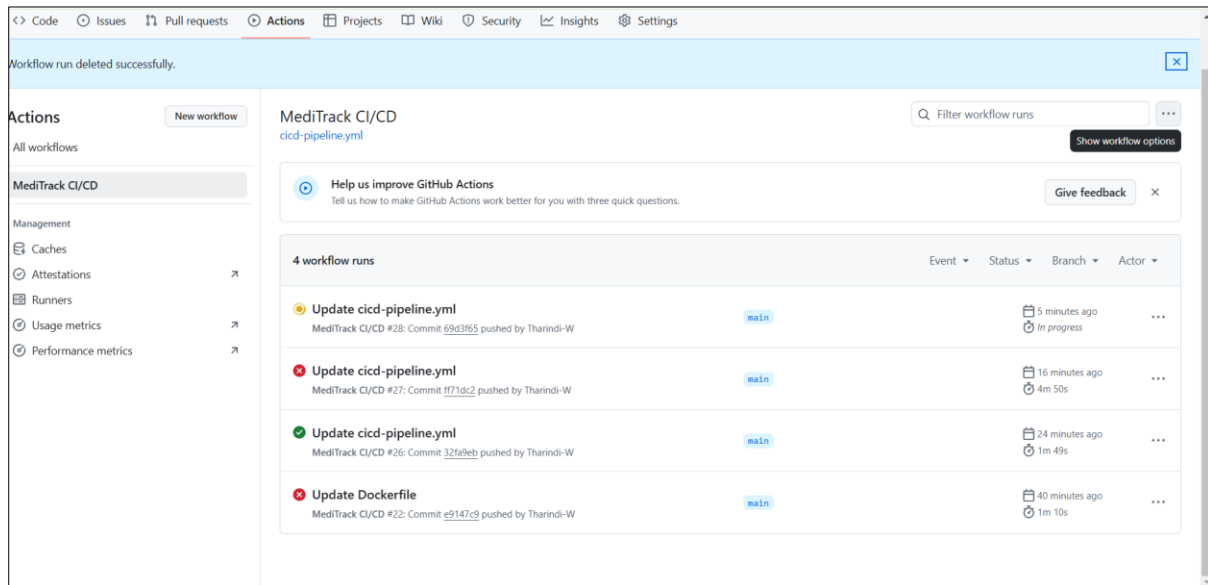


Figure 20: Github Actions Red-Green Workflow

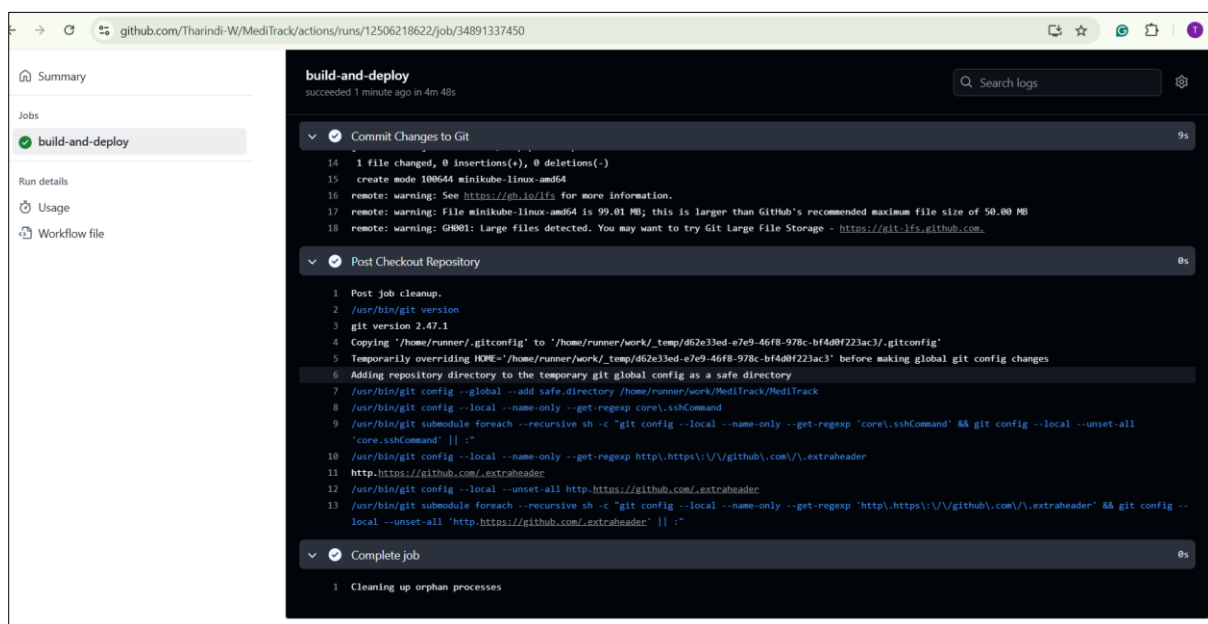


Figure 21: Github Action CI/CD

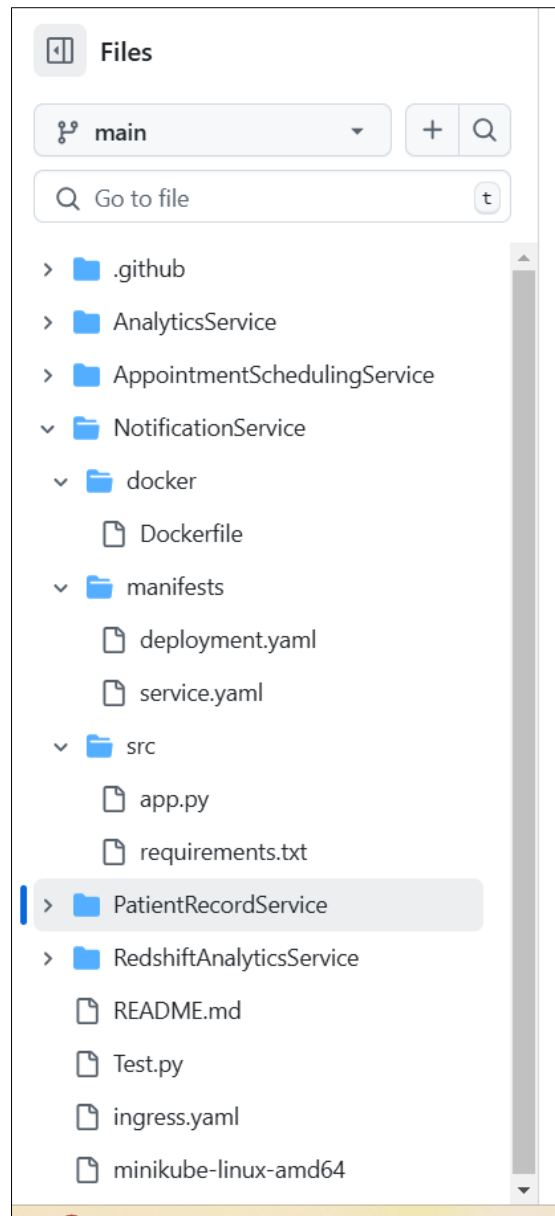


Figure 22: Github Folder Structure

9 SECURITY AND ETHICS

9.1.1 Security Challenges

1. Sensitive Data Protection (Patient Data),
Storing and transmitting patient health data without encryption exposes sensitive personal health information (PHI) to potential breaches. Therefore, implementing end to end encryption for data both in transit and at rest, utilizing industry standard encryption protocols like TLS for data transmission and AES for data storage or exploring the use of AWS Key Management Service (KMS) to manage and rotate encryption keys securely would help overcome this issue.
2. Access Control and Authentication
Exposed public repositories and improper handling of API keys and database credentials could lead to unauthorized access to the system. Implementing secure authentication mechanisms using OAuth or OpenID Connect for API access and microservice authentication. Storing sensitive information such as passwords, API keys, and database credentials in AWS Secrets Manager or similar secure services can be used. In the solution used GitHub Actions secrets for securely managing credentials in CI/CD pipelines.
3. Container Security
Running containers in the Minikube Kubernetes cluster without proper security measures exposes the system to vulnerabilities such as unauthorized access and malicious attacks. Regularly scanning container images for vulnerabilities using tools like Clair or Anchore. Implementing Pod Security Policies (PSP) and Role Based Access Control (RBAC) in Kubernetes to restrict access can be used.
4. Public Repository Risks
Exposing source code and configuration files in a public GitHub repository increases the risk of revealing sensitive information, such as API keys and database connection details. Avoided committing sensitive information by using .gitignore files and ensuring that all credentials are stored in environment variables or GitHub Actions secrets. Considering switching to private repositories for production code to better protect sensitive information can be used.

9.1.2 Ethics Challenges

1. Patient Consent and Data Transparency
Collecting, storing, and processing patient data without clear consent can raise concerns regarding patient autonomy and privacy. Implementing clear and transparent consent mechanisms during patient registration, ensuring patients are aware of what data is being collected and how it will be used has to be concerned.
2. Data Privacy and Compliance
Storing patient health data without ensuring compliance with relevant privacy regulations, such as HIPAA or GDPR, could lead to serious legal and ethical issues. Ensuring compliance with privacy regulations by implementing strong data access controls, audit logs, and user consent management systems. Adhering to data minimization practices by collecting only necessary patient information and implementing robust data protection measures must be implemented.

3. Fair Evaluation and Use of Analytics

The use of analytics and productivity metrics may lead to biased or unfair evaluations of healthcare providers or patients. Ensuring that any analytics or metrics used within the platform are employed responsibly and focused on improving patient care and provider performance. Providing transparency into how metrics are calculated and used to avoid misuse for unfair comparisons or evaluations is the best.

10 RUNBOOK

10.1 INSTALLING MINIKUBE

1. **Install Docker** - Minikube requires Docker as a container runtime. Install Docker using the following commands,

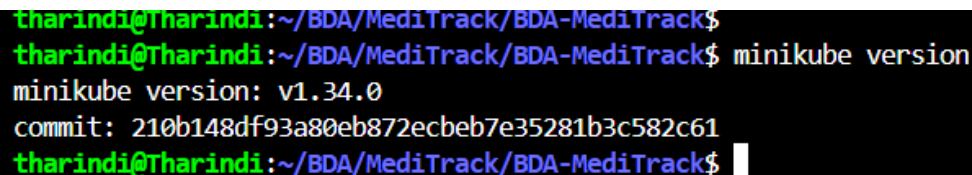
```
sudo apt update
sudo apt install docker.io -y
```

2. **Install Minikube** - Download and install Minikube from the official website or use the following commands,

```
curl -LO https://storage.googleapis.com/minikube/releases/latest/minikube-linux-amd64
sudo install minikube-linux-amd64 /usr/local/bin/minikube
```

3. **Verify Installation** - Run the command to confirm Minikube is installed,

```
minikube version
```

A terminal window with a black background and green text. The prompt is 'tharindi@Tharindi:~/BDA/MediTrack/BDA-MediTrack\$'. The command 'minikube version' has been entered and executed. The output shows 'minikube version: v1.34.0' and 'commit: 210b148df93a80eb872ecbeb7e35281b3c582c61'.

```
tharindi@Tharindi:~/BDA/MediTrack/BDA-MediTrack$
tharindi@Tharindi:~/BDA/MediTrack/BDA-MediTrack$ minikube version
minikube version: v1.34.0
commit: 210b148df93a80eb872ecbeb7e35281b3c582c61
tharindi@Tharindi:~/BDA/MediTrack/BDA-MediTrack$
```

Figure 23: Minikube Version Check

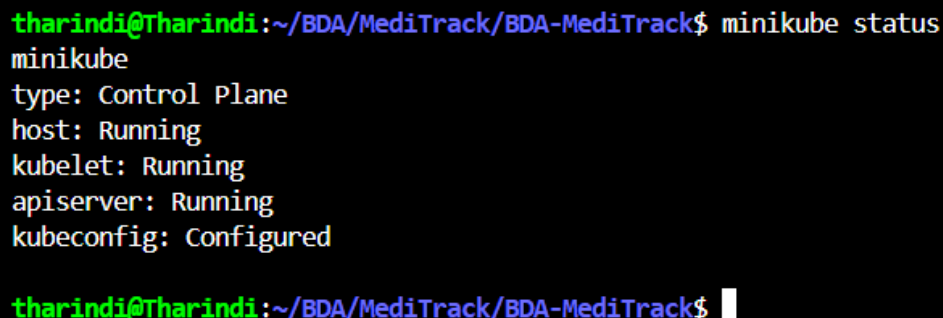
10.2 STARTING MINIKUBE

1. **Start Minikube** - Use the following command to start Minikube,

```
minikube start --driver=docker
```

2. **Verify Minikube Status** - Confirm Minikube is running,

```
minikube status
```

A terminal window with a black background and green text. The prompt is 'tharindi@Tharindi:~/BDA/MediTrack/BDA-MediTrack\$'. The command 'minikube status' has been entered and executed. The output shows 'minikube type: Control Plane', 'host: Running', 'kubelet: Running', 'apiserver: Running', and 'kubeconfig: Configured'.

```
tharindi@Tharindi:~/BDA/MediTrack/BDA-MediTrack$ minikube status
minikube
type: Control Plane
host: Running
kubelet: Running
apiserver: Running
kubeconfig: Configured
tharindi@Tharindi:~/BDA/MediTrack/BDA-MediTrack$
```

Figure 24: Minikube Status Check

10.3 ENABLING MINIKUBE ADD-ONS

1. **Enable Ingress** - Minikube's ingress addon is required for managing traffic,

```
minikube addons enable ingress
```

2. **Enable Dashboard** - To visually monitor Minikube resources,

```
minikube addons enable dashboard
```

3. **Access the Dashboard** - Open the Kubernetes dashboard,

```
minikube dashboard
```

```
tharindi@Tharindi:~/BDA/MediTrack/BDA-MediTrack$ minikube addons list
```

ADDON NAME	PROFILE	STATUS	MAINTAINER
ambassador	minikube	disabled	3rd party (Ambassador)
auto-pause	minikube	disabled	minikube
cloud-spanner	minikube	disabled	Google
csi-hostpath-driver	minikube	disabled	Kubernetes
dashboard	minikube	enabled ✓	Kubernetes
default-storageclass	minikube	enabled ✓	Kubernetes
efk	minikube	disabled	3rd party (Elastic)
freshpod	minikube	disabled	Google
gcp-auth	minikube	disabled	Google
gvisor	minikube	disabled	minikube
headlamp	minikube	disabled	3rd party (kinvolk.io)
helm-tiller	minikube	disabled	3rd party (Helm)
inaccel	minikube	disabled	3rd party (InAccel [info@inaccel.com])
ingress	minikube	enabled ✓	Kubernetes
ingress-dns	minikube	enabled ✓	minikube
inspektor-gadget	minikube	disabled	3rd party (inspektor-gadget.io)
istio	minikube	disabled	3rd party (Istio)

Figure 25: Minikube Addon list

10.4 SETTING UP MICROSERVICES

1. **Build Docker Images** - Create Docker images for your Flask and analytics microservices.

```
docker build -t patient-record-service:latest -f docker/Dockerfile .
```

2. **Push Images to Minikube** - Transfer the images to Minikube's internal Docker registry,

```
eval $(minikube docker-env)
docker images
```

3. **Deploy Microservices**

- Write Kubernetes manifests (YAML files) for Flask and analytics services.
- Apply the manifests,

```
kubectl apply -f analytics-service.yaml
```

Refer Figure 11

10.5 CONNECTING MONGODB

1. **Deploy MongoDB** – Use MongoDB Atlas and verify connectivity and networks
2. **Update Microservice Configuration**
 - Add the MongoDB connection URL to your Flask and analytics microservices' environment variables.
 - Redeploy the microservices with updated configurations.

```
AnalyticsService > manifests > ! deployment.yaml
5 spec:
10   template:
14     spec:
16       - name: analytics-service
19         ports:
20         - containerPort: 5003
21           protocol: TCP
22         env:
23         - name: MONGO_URI
24           value: "mongodb+srv://e17299:gFH00zihG8Pjj0pX@meditrack.wg7uo.mongodb.net/?retryWrites=true&w=majority&appName=Meditrack"
25
```

Figure 26: MongoDB URL

10.6 CONNECTING TO AMAZON REDSHIFT

1. **Configure Redshift**
 - Create a Redshift cluster in AWS and note its connection details (IAM roles, VPC, Subnets, Security groups, port, database name, etc.).
2. **Update Microservices**
 - Add the Redshift credentials to the analytics microservice configuration.
 - Install a Redshift Python driver (e.g., psycopg2) in the microservice using requirements.txt

```
RedshiftAnalyticsService > src > requirements.txt
1 flask
2 psycopg2-binary
3 pymongo
```

Figure 27: Redshift Python Driver in Requirements.txt

10.7 SETTING UP QUICKSIGHT

1. **Connect to Redshift,**
 - In the AWS QuickSight interface, add Redshift as a data source.
2. **Create Visualizations,**

- Use the connected dataset to build dashboards and reports.

Refer Figure 18 and Figure 19

10.8 CONFIGURING GITHUB ACTIONS

1. Set Up Workflow

- Add a `.github/workflows/deploy.yml` file with the following:

```
name: Deploy to Minikube
on:
  push:
    branches:
      - main
jobs:
  deploy:
    runs-on: ubuntu-latest
    steps:
      - name: Checkout code
        uses: actions/checkout@v3
      - name: Set up Docker
        uses: docker/setup-buildx-action@v2
      - name: Build and push Docker image
        run: |
          docker build -t flask-service:latest ./flask-service
          docker build -t analytics-service:latest ./analytics-service
      - name: Deploy to Minikube
        run: |
          kubectl apply -f flask-service.yaml
          kubectl apply -f analytics-service.yaml
```

10.9 TESTING THE WORKFLOW

1. Send Sample Requests,

- Use tools like curl or Postman to test the Flask microservice

```
curl -X POST http://<minikube_ip>:<node_port>
```

2. Verify Data Flow:

- Check MongoDB for data storage.
- Confirm analytics results in Redshift.

This runbook provides a step-by-step guide to deploying and testing a Minikube-based microservices solution, ensuring connectivity with MongoDB, Redshift, and QuickSight, with automation powered by GitHub Actions.

REFERENCE

<https://dev.to/bravinsimiyu/building-and-deploying-a-flask-application-to-kubernetes-cluster-58e9>

<https://medium.com/@ikbenezer/how-to-create-an-amazon-elastic-kubernetes-service-cluster-using-terraform-c615eca68c54>

<https://minikube.sigs.k8s.io/docs/>

<https://github.blog/enterprise-software/ci-cd/build-ci-cd-pipeline-github-actions-four-steps/>